

15-418/618, Spring 2026

Assignment 4

Parallel VLSI Wire Routing via MPI

Assigned:	Wed., Feb. 25
Due:	Wed., Mar. 18, 11:59 pm
Last day to handin:	Sat., Mar. 21

1 Overview

Building upon the previous assignment, you will now be using *MPI* to write a *message passing* version of the same application. Before you begin, please take the time to review the course policy on academic integrity at:

<http://www.cs.cmu.edu/~418/academicintegrity.html>

Download the Assignment 4 starter code from the course Github using:

```
git clone https://github.com/cmu15418s26/asst4.git
```

2 MPI Resources

If you do not have an experience with MPI, we strongly recommend going through the MPI tutorial at <https://hpc-tutorials.llnl.gov/mpi/> before you attempt doing the assignment. Learn how to compile and run MPI binaries on the same host and across the nodes of a cluster. Study, compile, and run at least a few of the exercises (Hello World, Array Decomposition, etc.) listed at https://hpc-tutorials.llnl.gov/mpi/exercise_1/. Other MPI resources include the documentation pages on [MPI Forum](#) and [Open MPI](#).

To use OpenMPI on the PSC machines, you must first run “`module load openmpi`” to load the OpenMPI module. (This loads “`openmpi/4.1.1-gcc8.3.1`” as the default, but you can also choose different compiler versions from the list after running “`module avail openmpi`”.) To run a program with MPI, use the `mpirun` command. The handout also provides an example MPI program `sqrt3` that approximates `sqrt(3)`.

We encourage you to use TAU Commander, an MPI profiling tool, that is pre-installed on the GHC machines at `/opt/taucmdr/bin`. You should add TAU Commander to your `PATH` environment variable with the

command: `export PATH=/opt/taucmdr/bin:$PATH`. You can confirm that the tool is installed correctly by running the command: `tau initialize` which should exit with status code 0. For more detailed instructions on how to use TAU Commander and visualize its output, we recommend the following guide: <https://hpc.llnl.gov/software/development-environment-software/tau-tuning-and-analysis-utilities>

3 Programming Task: Message-Passing Version of VLSI Wire Routing

Please refer to the Assignment 3 handout for a general discussion of the VLSI wire routing problem. Unless we specify otherwise, please use the same problem setup from Assignment 3 (e.g., wire routing constraints, structure of the simulated annealing algorithm, etc.).

Recall that in Assignment 3, we discussed two different parallelization strategies for the VLSI wire-routing problem: *within-wires* and *across-wires*. For this assignment, you are required to implement one of two options: either (i) the *across-wires* approach by itself, or (ii) a hybrid approach that combines both *across-wires* and *within-wires* together. You may not implement only the *within-wires* approach, however.

Keep in mind that because sending messages is relatively expensive, experienced MPI programmers generally try to structure their code to avoid frequent communication. We are interested in seeing good performance from your code, so if your initial approach does not scale well, be sure to explore various options for improving your code.

3.1 Batching Revisited

Recall that in Assignment 3, your *across-wires* strategy included a *batch size* parameter. In this assignment, you will also be using batching, although there are some differences in how this affects the orchestration step, given the differences between the shared-address-space and message-passing programming models. For the message-passing version, batching does not result in delays to updating the local (private) memory with the new route (since no other processors can access a thread's local memory), but instead it creates an opportunity to batch together updates into larger messages. An illustration of this batching approach is shown in Figure 1.

Similar to Assignment 3, the batch size parameter provides a mechanism to make tradeoffs between performance and the quality of the result. Increasing the batch size helps to amortize communication overheads, but it may also have a negative impact on the overall quality of the routing result (i.e. the overall cost may increase). Hence we would like for you to experiment with this batch size parameter (up to batch size = 16), to see how these tradeoffs affect both quality and performance. Specifically, we would like you compare batch sizes of 1, 4, and 16.

3.2 Optimizations that are prohibited.

Similar to Assignment 3, you may not use memoization tricks to reduce the amount of computation in the workload. While there are clever ways that one could modify the exhaustive search algorithm to dramatically reduce the amount of computation and memory accesses involved, these changes are not allowed on this assignment because they dramatically alter the arithmetic intensity and scalability properties of the

```

// Outer simulated annealing time steps
for (timestep = 0; timestep < N_iters; timestep++) {
    while (there is still work to do in this timestep) {
        each processor grabs a batch of B wires to be routed;
        for (i = 0; i < B; i++) {
            determine the new route for wire i (within the batch);
            if (the route has changed) {
                update local (private) memory with the new route;
                remember this updated route (to propagate it later);
            }
        }
        // Note: this propagation step can be done asynchronously
        propagate updated routes from this batch to other processors;
        // Remember to also process updates from other processors
    }
}

```

Figure 1: Pseudo-code for the outer control structure of a message-passing version of the Across-Wires approach with batching (where B is the batch size).

computation. For example, you must NOT do any of the following: (i) sample a subset of the potential paths (or short-circuit the exploration of paths); (ii) use prefix sums; or (iii) use range-query data structures or algorithms. The point of the assignment is to explore ways to parallelize the given workload.

4 General Advice

Implications of private address spaces: This is our first assignment where you only have private address spaces. This requires a change of thinking when you design your algorithm, and it is important to wrap your head around the implications of having only private memory. For example, do not be tempted to try to mimic a shared address space by either sending frequent broadcast messages or treating a single processor as the the memory module, since these approaches do not scale well. Go back and review the message-passing version of the Grid Solver program (in Lecture 5), and think about how it differed from the other versions of the Grid Solver.

Be careful about using MPI collective communication primitives: MPI includes a set of [primitives for collective communication](#), including broadcast (`MPI_Bcast`). While these primitives may be convenient for programmers, it is important to realize that the underlying implementations of these primitives are usually quite expensive. Hence experienced MPI programmers use these primitives sparingly: e.g., to distribute input parameters to all of the processors at the start of the program, to communicate scalar variables at time step boundaries (as we saw with the Grid Solver example in Lecture 5), etc. If you find yourself tempted to use the collective communication primitives frequently to communicate large amounts of data, be forewarned that this may become a significant performance bottleneck for your program (especially at higher thread counts). It is okay to use these operations judiciously, especially if you perform measurements to confirm that they are not hurting scalability.

Every processor counts: Sometimes novice MPI programmers are tempted to set aside one processor (e.g., processor 0) as a “manager” that distributes work to all of the other processors (the “workers”). While this approach may simplify some aspects of writing the code, it has some unfortunate scalability problems. First, when we calculate speedup when running on P processors, this includes *all* of the P processors, not just the $P - 1$ workers. Therefore if you set one processor aside as a manager and have only $P - 1$ processors doing parallel work, it immediately hurts your scalability (especially with smaller thread counts). A second performance problem with this manager/worker approach occurs with higher thread counts: communication with the manager tends to become a bottleneck in the interconnection network, since the $P - 1$ workers are all trying to send and receive messages from it at the same time (causing it to become a hot spot, often incurring significant queueing delays).

Hence experienced MPI programmers prefer to use all P processors to perform parallel computation in a fairly symmetric fashion with peer-to-peer communication, avoiding centralized coordination when possible. Many different communication structures can achieve this. For example, one option is a software-managed ring, which enables processors to use point-to-point send/receive messages to communicate effectively. In a ring, when a processor receives a message that it originally initiated, it knows the message has been seen by all other processors. However, rings are just one possible structure. Other symmetric communication patterns (grid, mesh, etc) may be more appropriate depending on the algorithm and workload.

Think carefully about the design of your messages: In the Grid Solver example from Lecture 5, most of the messages that were sent were simply boundary rows. For the assignment, you may need to put more thought into the structure of your messages. There are opportunities for reducing communication overhead if you think about this carefully. At the same time, be careful about the functionality of propagating updates asynchronously (you don’t want to end up with corrupted occupancy arrays).

Differences between sending messages within and across chips: Within a processor chip, message passing can be implemented (under the covers) using memory copy operations. Across separate processor chips, however, messages are passed across interconnection networks. Therefore the latency and bandwidth properties of sending messages in MPI can change noticeably once the computation scales beyond individual processor chip boundaries (i.e. at higher thread counts).

5 Logistical Details

Unless specified otherwise, please refer to the details in the Assignment 3 handout regarding command line flags, the formats of the input and output files, and strategies for measuring performance.

Input files: The set of input files will remain the same as Assignment 3 (located in `/code/inputs`).

Guidelines on performance and cost targets: To give you a sense of whether your code is in the right ballpark in terms of its performance and the quality of the routing results (i.e. the cost metrics), we will be posting guidelines on these numbers via Piazza by February 25th.

Using the GHC and PSC machines. Compared with Assignment 3, we will be doing more extensive performance analysis on the PSC machines for this assignment since they have more cores and faster

interconnection networks. However, it is still the case that you should start on the GHC machines to do all of your initial debugging and testing of your code. (Host names for these machines are `ghcX.ghc.andrew.cmu.edu`, where X is between 26 and 86.) Once your code is stable and scaling well on the GHC machines, please begin running experiments on the PSC Bridges-2 Regular Memory (RM) machines. Here is a [tutorial on accessing the PSC machines](#). You may also find it helpful to refer to the [Bridges-2 User Guide](#). **IMPORTANT:** To help us preserve our limited allocation of time on the PSC machines, please follow these rules regarding using the PSC machines:

- Do not perform any debugging (or significant editing) on the PSC machines. We are charged the same for debugging time as we are for collecting performance data, and this time is precious. You should debug on GHC, and only transfer code to the PSC machines when it is stable.
- Collect a full set of performance results on the GHC machines before you start running experiments on the PSC machines. Once you are happy with your performance debugging on GHC, then start doing your PSC experiments.
- Submit no more than 30 sets of experiments per day on the PSC machines. (A “set” of experiments is the data necessary for one speedup curve: i.e. a particular configuration on 1, 2, 4, 8, 16, 32, 64, and 128 threads.) There is a significant lag between when machine time is used up and when the instructors are notified about it. In previous years, we have had problems where a few groups performed large numbers of experiments, using up the bulk of the time allocation for the class (resulting in their PSC accounts being suspended). So please be careful to stay within your 30-experiments-per-day budget. Thanks!

6 Performance and Cost Targets

When we evaluate the performance of your code, we care about three metrics: (i) scalability (i.e. how much speedup you achieve as you scale up to more processors), (ii) absolute performance, and (iii) the overall cost metric for your wire routing result.

We do not provide a single target metric because there are ways to artificially boost some of these three metrics by penalizing the others. For example, if the processors did not communicate with each other at all during wiring, they might achieve excellent speedup at the expense of extremely poor wiring cost. As other examples, it is sometimes possible to achieve excellent speedup by artificially slowing down single-thread performance, or to hit an absolute performance target on 8 threads by focusing only on boosting single-thread performance on non-scalable code. We definitely do not want you to play games like this.

Here is the way to think about what we are looking for: assuming that your code has reasonably good single-thread absolute performance and produces a wiring result with reasonably good cost, we are especially interested in the *scalability* of your code (i.e. how well it speeds up as you increase the number of processors). (Note: you are likely to find scalability a bit more challenging in this assignment than it was in assignment 3.)

To give you a more concrete sense of what we are looking for, we attached performance benchmarks to the end of this file. When we assign your overall score on the assignment, 30% of the score will be based upon the performance of your code, and 70% will be based upon the quality of your discussion and analysis in your writeup. Our performance guidance will include two target numbers: one to achieve 50% of the

performance points, and one to achieve 80% of the performance points. To achieve credit beyond 80% on the performance score, we will be looking at the distribution of the performance of the top performers in the class. (Note that once you achieve the 80% target, the course staff will not be providing advice on how to optimize your code further: that is up to your group to explore on your own.)

7 Your Mission

In this assignment, you will be implementing, evaluating, and iteratively improving a message-passing version of the wire routing problem using the MPI programming model. The goal of this assignment is for you to think carefully about how real-world effects in the machine are limiting your speedup, and how you can improve your program to get better performance. If your performance is disappointing, then it is likely that you can restructure your code to make things better.

For this assignment, we require you to implement the *across-wires* approach or a hybrid approach. You cannot implement only the *within-wires* approach.

In your write-up, please include the following items:

1. **[40 points] Design and performance debugging journey for your message-passing approach:** Provide a detailed discussion of the thought process that went into designing your program, and how it evolved over time based upon your experiments. (See the Piazza post entitled “[Performance Debugging](#)”). Specifically, try to address the following questions:
 - What approaches did you take to parallelize the algorithm (including not only your final design, but any other designs along the way)?
 - Include your reasoning for your final implementation choices, including any graphs or tables that helped you make your decisions.
 - Where is the synchronization in your solution? Did you do anything to limit the overhead of synchronization?
 - Why do you think your code is unable to achieve perfect speedup? (Is it workload imbalance? communication/synchronization? data movement?)
 - At high thread counts, do you observe a drop-off in performance? If so (and you may not), why do you think this might be the case?
2. **[15 points] Experimental results from the GHC machines:** For each of these experiments, please collect and present data for 1, 2, 4, and 8 threads for each experiment while running on the GHC cluster (`ghcX`, where X is between 26 and 86).
 - (a) **Speedup graphs:** Show a plot of the *Total Speedup* and *Computation Speedup* vs. *Number of Processors* (N_{procs}). Discuss these results, including any non-ideal behaviors.
 - (b) **Cache misses:**
 - i. **Total cache misses:** Show a plot of the total number of *cache misses* for the entire program vs. *Number of Processors* (N_{procs}).

- ii. **Per-thread cache misses:** Show plot of the arithmetic mean of *per-thread* cache misses (from `perf stat -e cache-misses $PROGRAM`) vs. *Number of Processors* (N_{procs}).
 - iii. **Discussion:** Discuss the trends that you see in these cache miss plots, including any surprises, and how these numbers relate to your speedup graphs.
 - (c) **Sensitivity to batch size:** Show plots of how varying the per-processor *batch size* (B) between 1, 4, and 16 affects both performance and the overall cost metric. Please run these experiments on 8 threads. Discuss the impact of varying B on performance and overall cost, explaining any effects that you see.
3. [20 points] **Experimental results from the PSC machines:** Measure both the performance and the quality of the output for your parallel code running on the PSC Bridges-2 Regular Memory machines. Please collect and present data for 1, 2, 4, 8, 16, 32, 64, and 128 threads.
- (a) **Speedup graphs:** Show a plot of the *Total Speedup* and *Computation Speedup* vs. *Number of Processors* (N_{procs}). Discuss these results, including any non-ideal behaviors.
 - (b) **Comparison with GHC results:** Discuss how these results compare with the speedup curves that you measured on the GHC machines.
 - (c) **Sensitivity to batch size:** Show plots of how varying the per-processor *batch size* (B) between 1 and 16 affects both performance and the overall cost metric. Please run these experiments on 32 threads. Discuss the impact of varying B on performance and overall cost, explaining any effects that you see.
 - (d) **Scaling across multiple nodes:** Evaluate the performance of scaling beyond a single node to multiple nodes. Please collect and present data for 2, 8, and 64 threads on a single node versus the same number of threads (per node) running on 2 nodes (i.e. a total of 4, 16, and 128 threads in the latter case). Discuss your results, and any lessons that you have learned from your analysis.

8 Hand-in Instructions

You will submit your code via Autolab and your report via Gradescope. For the code, you will be submitting all C++ header and source files in the `code` folder.

1. Submitting your code:
 - (a) If you are working with a partner, form a group on Autolab. Do this before submitting your assignment. One submission per group is sufficient.
 - (b) Make sure all of your code is compilable and runnable. We should be able to simply run `make`. Please remove excessive print statements, if they were added.
 - (c) Run the command “`make handin.tar`.” This will run “`make clean`” and then create an archive of any C++ source code in `code`. If you find it absolutely necessary, you may modify the Makefile to include other files you have added.
 - (d) Submit the file `handin.tar` to Autolab.

2. Submitting your writeup:

- (a) Please upload your report as file report.pdf to Gradescope, one submission per team, and select the appropriate pages for each part of the assignment. After submitting, you will be able to add your teammate using the add group members button on the top right of your submission.

9 Benchmark of A Good Performance Solution

With $n = 1$, Table 2 reports the computation time on GHC and PSC machines. The hyper-parameters used here are: SA_prob=0.1, SA_iters=5, B=1.


Table 1: Computation time (seconds) with $n = 1$ on GHC and PSC.

File Name	Num Wires	Cost	t (GHC)	t (PSC)
few_wires.txt	128	29,240	2.543	3.400
medium_wires.txt	2048	267,589	8.485	11.134
abundant_wires.txt	16384	1,539,198	20.709	26.779

Table 2: Cost for medium_wires.txt on GHC machines.


Costs	1 Thread	2 Threads	4 Threads	8 Threads
batch = 1	267179	267261	267517	268317
batch = 4	267641	267547	268053	267959
batch = 16	267587	267555	267639	268051

Speedup on GHC Machines (batch 4)



(a) Speedup on GHC machines


Speedup on PSC Machines (batch 4)



(b) Speedup on PSC machines


Figure 2: Speedup for batch size 4.

Speedup on GHC Machines (batch 16)



(a) Speedup on GHC machines

Speedup on PSC Machines (batch 16)



(b) Speedup on PSC machines

Figure 3: Speedup for batch size 16.